

# Final Review Session

May 8, 2024

## 0.0.1 Final Review Session

```
[3]: from sklearn.preprocessing import LabelBinarizer
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

### Handling Missing Values

```
[4]: data = {
    'Employee_ID': [101, 102, 103, 104, 105, 106, 107, 108, 109, 110],
    'Name': ['Ahmed', 'Bob', 'Charlie', 'David', 'Eve', 'Corey', 'Hussain', 'Jacob', 'Motaz', 'Elmer'],
    'Age': [25, np.nan, 30, 35, 40, 22, 43, 32, 21, 65],
    'Department': ['HR', 'IT', np.nan, 'Finance', 'Operations', 'Finance', 'Operations', 'HR', 'IT', 'IT'],
    'Salary': [50000, 60000, 70000, np.nan, 80000, 50000, 60000, 17000, 20000, 40000]
}

emp_df = pd.DataFrame(data)
emp_df
```

```
[4]:
```

	Employee_ID	Name	Age	Department	Salary
0	101	Ahmed	25.0	HR	50000.0
1	102	Bob	NaN	IT	60000.0
2	103	Charlie	30.0	NaN	70000.0
3	104	David	35.0	Finance	NaN
4	105	Eve	40.0	Operations	80000.0
5	106	Corey	22.0	Finance	50000.0
6	107	Hussain	43.0	Operations	60000.0
7	108	Jacob	32.0	HR	17000.0
8	109	Motaz	21.0	IT	20000.0
9	110	Elmer	65.0	IT	40000.0

```
[40]: emp_df[np.sum(emp_df.isna(), axis=1)==0]
```

```
[40]: Employee_ID    Name    Age  Department    Salary
      0         101    Ahmed  25.0         HR  50000.0
      4         105      Eve  40.0    Operations  80000.0
      5         106    Corey  22.0         Finance  50000.0
      6         107  Hussain  43.0    Operations  60000.0
      7         108    Jacob  32.0         HR  17000.0
      8         109    Motaz  21.0         IT  20000.0
      9         110    Elmer  65.0         IT  40000.0
```

Axis 0 (column wise) vs Axis 1 (row wise)

```
[32]: np.sum(emp_df.isna(), axis=0)
```

```
[32]: Employee_ID    0
      Name          0
      Age          1
      Department    1
      Salary        1
      dtype: int64
```

```
[41]: np.sum(emp_df.isna(), axis=1)
```

```
[41]: 0    0
      1    1
      2    1
      3    1
      4    0
      5    0
      6    0
      7    0
      8    0
      9    0
      dtype: int64
```

Group By with Aggregation

```
[6]: emp_df[['Age', 'Department', 'Salary']].groupby('Department').agg({'Age': 'mean', 'Salary': ('min', 'max')})
```

```
[6]:           Age    Salary
      mean      min      max
Department
Finance    28.5  50000.0  50000.0
HR         28.5  17000.0  50000.0
IT         43.0  20000.0  60000.0
Operations  41.5  60000.0  80000.0
```

### Example: Library Data

```
[7]: data = {
    'ID': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'Floor': [1, 1, 2, 2, 3, 3, 4, 4, 5, 5],
    'Section Name': ['Fiction', 'Non-Fiction', 'Science', 'Mathematics', 'History',
                    'Biography', 'Poetry', 'Art', 'Children', 'Reference'],
    'Section Head': ['John Doe', 'John Doe', 'Jane Smith', 'Jane Smith', 'Mark Johnson',
                    'Mark Johnson', 'Emily Brown', 'Emily Brown', 'Sarah Wilson', 'Sarah Wilson'],
    'No. of Books Available': [100, 80, 120, 90, 110, 70, 95, 75, 85, 100],
    'No. of Books Lent': [20, 15, 25, 10, 30, 5, 15, 10, 10, 20]
}

lib_df = pd.DataFrame(data)
lib_df
```

```
[7]:
```

	ID	Floor	Section Name	Section Head	No. of Books Available \
0	1	1	Fiction	John Doe	100
1	2	1	Non-Fiction	John Doe	80
2	3	2	Science	Jane Smith	120
3	4	2	Mathematics	Jane Smith	90
4	5	3	History	Mark Johnson	110
5	6	3	Biography	Mark Johnson	70
6	7	4	Poetry	Emily Brown	95
7	8	4	Art	Emily Brown	75
8	9	5	Children	Sarah Wilson	85
9	10	5	Reference	Sarah Wilson	100

	No. of Books Lent
0	20
1	15
2	25
3	10
4	30
5	5
6	15
7	10
8	10
9	20

```
[8]: lib_df[['Section Name', 'Section Head', 'No. of Books Available', 'No. of Books Lent']].groupby(['Section Head', 'Section Name']).sum()
```

```
[8]:
```

		No. of Books Available	No. of Books Lent
Section Head	Section Name		

Emily Brown	Art	75	10
	Poetry	95	15
Jane Smith	Mathematics	90	10
	Science	120	25
John Doe	Fiction	100	20
	Non-Fiction	80	15
Mark Johnson	Biography	70	5
	History	110	30
Sarah Wilson	Children	85	10
	Reference	100	20

### Data Filtering

```
[26]: (lib_df['No. of Books Lent']>= 20) & (lib_df['No. of Books Available']< 115)
```

```
[26]: 0    True
      1   False
      2   False
      3   False
      4    True
      5   False
      6   False
      7   False
      8   False
      9    True
      dtype: bool
```

```
[27]: lib_df[(lib_df['No. of Books Lent']>= 20) & (lib_df['No. of Books Available']< 115)]
```

```
[27]:   ID  Floor Section Name  Section Head  No. of Books Available \
0    1     1     Fiction      John Doe             100
4    5     3     History    Mark Johnson             110
9   10     5    Reference    Sarah Wilson             100

      No. of Books Lent
0                20
4                30
9                20
```

### Label Binarization

```
[11]: data = ['cat', 'dog', 'bird', 'cat', 'dog']
      lb = LabelBinarizer()
      binary_data = lb.fit_transform(data)
      print(binary_data)
```

```
[[0 1 0]
```

```
[0 0 1]
[1 0 0]
[0 1 0]
[0 0 1]]
```

```
[14]: data = {
        'ID': [1, 2, 3, 4, 5],
        'Name': ['Chair', 'Table', 'Sofa', 'Bed', 'Desk'],
        'Type': ['Wooden', 'Metal', 'Plastic', 'Wooden', 'Metal']
    }

    # Convert data into DataFrame
    df = pd.DataFrame(data)

    # Perform label binarization on the 'Type' column
    lb = LabelBinarizer()
    binary_type = lb.fit_transform(df['Type'])

    # Convert binary_type into DataFrame
    binary_type_df = pd.DataFrame(binary_type, columns=lb.classes_)

    # Concatenate the binary_type_df with the original DataFrame
    df = pd.concat([df, binary_type_df], axis=1)

    df
```

```
[14]:
```

	ID	Name	Type	Metal	Plastic	Wooden
0	1	Chair	Wooden	0	0	1
1	2	Table	Metal	1	0	0
2	3	Sofa	Plastic	0	1	0
3	4	Bed	Wooden	0	0	1
4	5	Desk	Metal	1	0	0

## Binning

```
[12]: import pandas as pd

data = {'Age': [22, 35, 47, 55, 68, 72, 28, 32, 45, 51]}

df = pd.DataFrame(data)

df
```

```
[12]:
```

	Age
0	22
1	35
2	47
3	55
4	68

```

5    72
6    28
7    32
8    45
9    51

```

```

[13]: # Define bin edges
bins = [0, 30, 50, 100]

labels = ['Young', 'Middle-aged', 'Senior']

df['Age Group'] = pd.cut(df['Age'], bins=bins, labels=labels)

df

```

```

[13]:   Age  Age Group
0    22    Young
1    35 Middle-aged
2    47 Middle-aged
3    55    Senior
4    68    Senior
5    72    Senior
6    28    Young
7    32 Middle-aged
8    45 Middle-aged
9    51    Senior

```

## Data Visualization Examples

```

[15]: import matplotlib.pyplot as plt

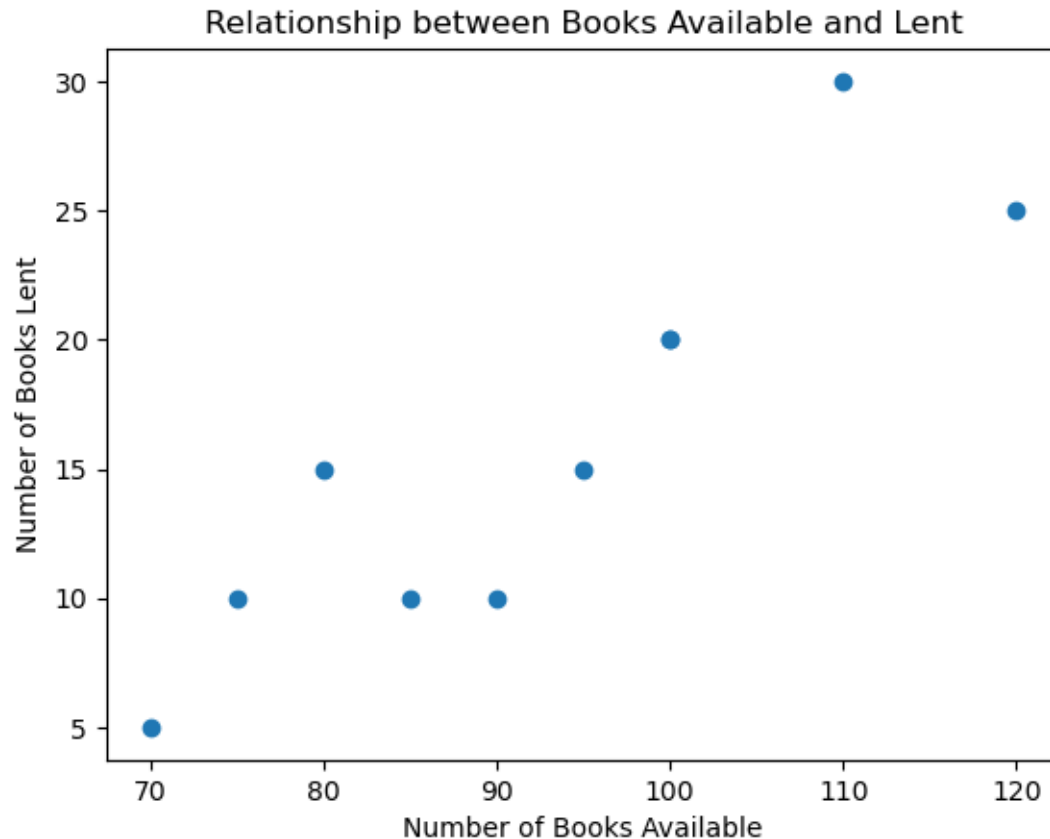
# Group the data by Section Name and calculate the sum of books available and
↳lent
section_stats = lib_df.groupby('Section Name').sum()

# Plot bar chart
section_stats.plot(kind='bar', y=['No. of Books Available', 'No. of Books_
↳Lent'],
                    xlabel='Section Name', ylabel='Number of Books',
                    title='Number of Books Available and Lent by Section')
plt.xticks(rotation=45)
plt.show()

```



```
[16]: # Plot scatter plot
plt.scatter(lib_df['No. of Books Available'], lib_df['No. of Books Lent'])
plt.xlabel('Number of Books Available')
plt.ylabel('Number of Books Lent')
plt.title('Relationship between Books Available and Lent')
plt.show()
```



```
[18]: np.random.seed(0)
np.random.seed(0)
time = pd.date_range(start='2024-01-01', periods=350, freq='H')
speed = np.random.randint(40, 80, size=350)
category = np.random.choice(['Fast', 'Slow'], size=350)

traffic_data = pd.DataFrame({'Time': time, 'Speed': speed, 'Category':
    category})
traffic_data.head()
```

/tmp/ipykernel\_2066/2314165996.py:3: FutureWarning: 'H' is deprecated and will be removed in a future version, please use 'h' instead.

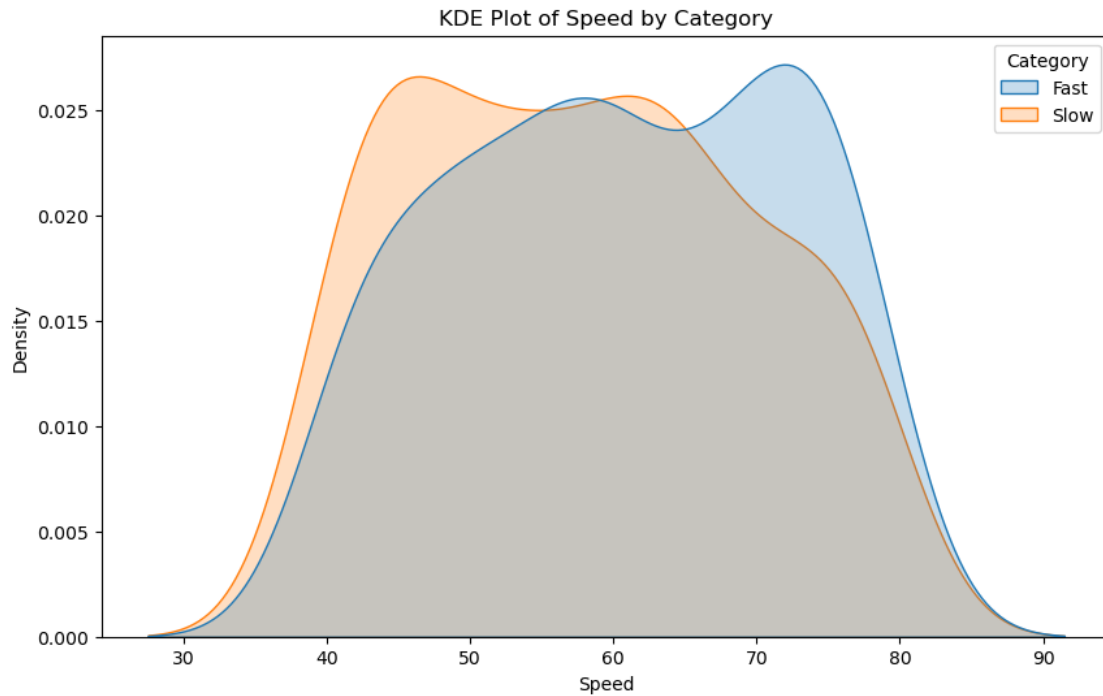
```
time = pd.date_range(start='2024-01-01', periods=350, freq='H')
```

```
[18]:
```


	Time	Speed	Category
0	2024-01-01 00:00:00	40	Fast
1	2024-01-01 01:00:00	43	Fast
2	2024-01-01 02:00:00	43	Slow
3	2024-01-01 03:00:00	79	Slow
4	2024-01-01 04:00:00	49	Fast



```
[19]: plt.figure(figsize=(10, 6))
sns.kdeplot(data=traffic_data, x='Speed', hue='Category', fill=True,
            common_norm=False)
plt.title('KDE Plot of Speed by Category')
plt.xlabel('Speed')
plt.ylabel('Density')
plt.show()
```



```
[20]: plt.figure(figsize=(10, 6))
sns.stripplot(data=traffic_data, x='Category', y='Speed', jitter=True)
plt.title('Strip Plot of Speed by Category')
plt.xlabel('Category')
plt.ylabel('Speed')
plt.show()
```



## Web Scraping

```
[21]: import requests
from bs4 import BeautifulSoup

url = 'https://www.geeksforgeeks.org/python-programming-language/'
response = requests.get(url)
soup = BeautifulSoup(response.text, 'html.parser')
x = soup.find_all('h2')
for z in x:
    print(z.text)
```

What is Python?

Writing your first Python Program to Learn Python Programming

Python3

Table of Content

Setting up Python

Getting Started with Python Programming

Learn Python Input/Output

Python Data Types

Python Operators

Python Conditional Statement

Python Loops

Python Functions

Python OOPs Concepts  
 Python Exception Handling  
 Python Packages or Libraries  
 Python Collections  
 Python Database Handling  
 Python vs. Other Programming Languages  
 Learn More About Python with Different Applications:  
 Python Online Quiz  
 Python Latest & Upcoming Features  
 What kind of Experience do you want to share?

```
[22]: url = 'https://www.geeksforgeeks.org/python-programming-language/'
response = requests.get(url)
soup = BeautifulSoup(response.text, 'html.parser')
x = soup.find_all('a')
for z in x[-5:]:
    print(z.text.strip())
```

Work Experiences  
 Campus Experiences  
 Competitive Exam Experiences  
 Can't choose a topic to write? [click here](#) for suggested topics  
 Write and publish your own Article

### Pivot vs. Group By

```
[33]: # importing pandas
import pandas as pd

# creating dataframe
df = pd.DataFrame({'Product': ['Carrots', 'Broccoli', 'Banana', 'Banana',
                                'Beans', 'Orange'],
                   'Category': ['Vegetable', 'Vegetable', 'Fruit', 'Fruit',
                                'Vegetable'],
                   'Quantity': [8, 5, 3, 4, 5, 9, 11, 8],
                   'Amount': [270, 239, 617, 384, 626, 610, 62, 90]})
df
```

```
[33]:
```

	Product	Category	Quantity	Amount
0	Carrots	Vegetable	8	270
1	Broccoli	Vegetable	5	239
2	Banana	Fruit	3	617
3	Banana	Fruit	4	384
4	Beans	Vegetable	5	626

5	Orange	Fruit	9	610
6	Broccoli	Vegetable	11	62
7	Banana	Fruit	8	90

```
[35]: pivot = df.pivot_table(index=['Product'],
                             values=['Amount'],
                             aggfunc='sum')
pivot
```

```
[35]:
```

	Amount
Product	
Banana	1091
Beans	626
Broccoli	301
Carrots	270
Orange	610

```
[37]: df[['Product', 'Amount']].groupby('Product').sum()
```

```
[37]:
```

	Amount
Product	
Banana	1091
Beans	626
Broccoli	301
Carrots	270
Orange	610

### Difference between Pivot and Group By:

- Pivot is primarily used for reshaping or restructuring data, rotating rows into columns or vice versa whereas Group by is primarily used for aggregation, summarizing data based on one or more key columns.
- Pivot produces a new table with reshaped data, often resulting in a multi-level index or hierarchical columns whereas Group by produces a summary or aggregation of data, typically in a Series or DataFrame with aggregated values.
- Pivot allows you to specify columns to use as index, columns, and values whereas Group by allows you to specify one or more key columns for grouping and apply aggregation functions on one or more columns.

```
[ ]:
```